

```
from google.colab import files
```

```
uploaded=files.upload()
```

```
import pandas as pd
```

```
df=pd.read_csv("Telco-Customer-Churn.csv")
```

```
df.head()
```

```
#Data Exploration
```

```
df.info()
```

```
df.describe()
```

```
df.columns
```

```
df.shape
```

```
#checking missing value and duplicates
```

```
print(df.isnull().sum())
```

```
print(f"Duplicated Rows :{df.duplicated().sum()}")
```

```
#visualize features
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
sns.countplot(x='Churn', data=df)
```

```
plt.title('Distribution of Churn')
```

```
plt.show()
sns.histplot(df['tenure'],kde=True)
plt.title('Distribution of Tenure')
plt.show()
sns.boxplot(x='Churn', y='MonthlyCharges', data=df)
plt.title('Monthly Charges vs Churn')
plt.show()
```

#identifying target and features

```
target='Churn'
```

```
features=df.drop(columns=[target]).columns.tolist()
```

#convert catgo to numeric

```
df['Totalcharges']=pd.to_numeric(df['TotalCharges'],errors='coerce') df.dropna(inplace=True)
```

#one-hot encode

```
df_encoded=pd.get_dummies(df,drop_first=True)
```

#featue scaling

```
from sklearn.preprocessing import StandardScaler
```

```
scaler=StandardScaler()
```

```
scaled_features=scaler.fit_transform(df_encoded.drop(columns=['Churn_Yes']))
```

```
X=pd.DataFrame(scaled_features,columns=df_encoded.drop(columns=['Churn_Yes']).columns) y=df_encoded['Churn_Yes']
```

#train-test split

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
2,random_state=42)
```

```
df['Churn'].value_counts()
```

```
# Apply SMOTE to balance the dataset
```

```
from imblearn.over_sampling import SMOTE smote =
SMOTE(random_state=42) X_resampled, y_resampled =
smote.fit_resample(X_train, y_train)
```

```
#model building
```

```
from sklearn.ensemble import RandomForestClassifier from
sklearn.linear_model import LogisticRegression from
sklearn.metrics import accuracy_score
```

```
# Random Forest (balanced)
```

```
rf_model = RandomForestClassifier(class_weight='balanced')
rf_model.fit(X_resampled, y_resampled)
```

```
# Logistic Regression (balanced)
```

```
lr_model = LogisticRegression(class_weight='balanced',
max_iter=1000) lr_model.fit(X_resampled, y_resampled)
```

```
#evaluation
```

```
from sklearn.metrics import classification_report,
confusion_matrix, accuracy_score
```

```
# Random Forest Evaluation
```

```
print("Random Forest Results:")
rf_model.fit(X_train, y_train)
rf_preds = rf_model.predict(X_test)
print(confusion_matrix(y_test, rf_preds))
print(classification_report(y_test, rf_preds))
print("RF Accuracy:", accuracy_score(y_test, rf_preds))

# Logistic Regression Evaluation
print("\nLogistic Regression Results:")
lr_model.fit(X_train, y_train)
lr_preds = lr_model.predict(X_test)
print(confusion_matrix(y_test, lr_preds))
print(classification_report(y_test, lr_preds))
print("LR Accuracy:", accuracy_score(y_test, lr_preds))

print(f"After SMOTE - Churn Distribution:
{y_resampled.value_counts()}")

# Compare accuracies and choose the better model
if accuracy_score(y_test, rf_preds) > accuracy_score(y_test,
lr_preds):
    model = rf_model
    print("Selected RF model")
else:
```

```
model = lr_model  
print("Selected LR model")
```

```
import joblib  
joblib.dump(model, 'churn_prediction_model.pkl')  
joblib.dump(X.columns.tolist(), 'columns.pkl')
```

```
!pip install gradio pandas joblib --quiet
```

```
import gradio as gr  
import pandas as pd import joblib  
model = joblib.load('churn_prediction_model.pkl') columns =  
joblib.load('columns.pkl')  
def predict_churn(gender, senior_citizen, partner,  
dependents, tenure, monthly, total, phone_service,  
multiple_lines, internet_service):  input_data = {  
    'gender': gender,  
    'SeniorCitizen': senior_citizen,  
    'partner': partner,  
    'dependents': dependents,  
    'tenure': tenure,  
    'MonthlyCharges': monthly,  
    'TotalCharges': total,  
    'PhoneService': phone_service,
```

```

'MultipleLines': multiple_lines,
    'InternetService': internet_service
}

df_input = pd.DataFrame([input_data])

df_encoded =
pd.get_dummies(df_input).reindex(columns=columns,
fill_value=0)  prediction = model.predict(df_encoded)

    return "Churn" if prediction[0] == 1 else "No Churn"

# Create Gradio interface

iface = gr.Interface (
    fn=predict_churn,
    inputs=[

        gr.Dropdown(['Female', 'Male'], label="Gender"),
        gr.Dropdown(['No', 'Yes'], label="Senior Citizen"),
        gr.Dropdown(['No', 'Yes'], label="Partner"),

        gr.Dropdown(['No', 'Yes'], label="Dependents"),

        gr.Slider(0, 72, label="Tenure (months)"),
        gr.Number(minimum=0.0, label="Monthly Charges"),
        gr.Number(minimum=0.0, label="Total Charges"),
        gr.Dropdown(['No', 'Yes'], label="Phone Service"),
        gr.Dropdown(['No phone service', 'No', 'Yes'], label="Multiple
Lines"),      gr.Dropdown(['DSL', 'Fiber optic', 'No'],
label="Internet Service")  ],  outputs="text",
    title="Customer Churn Prediction",
    description="Enter customer details to predict churn." )

```

```
# Launch the interface and share it publicly  
iface.launch(share=True)
```